

Network Layer

CS5700 Fall 2019

Where we are?

- Do you remember the responsibility of each layer?



Agenda

- Overview
- IP (v4, v6)
- Routing protocols
- ICMP

Overview

Two key network layer functions



- **Forwarding:** move packets from router's input port to appropriate output port
 - Local, per router function
- **Routing:** determine route taken by packets from source to destination
 - Network wide logic
 - Determine how datagram is routed among routers along end-to-end path from source to destination

Network service model

- Reliability?
- Order?
- Delay?
- Throughput?



Inside a router



Destination-based forwarding table

<i>forwarding table</i>	
Destination Address Range	Link Interface
11001000 00010111 00010000 00000000 through 11001000 00010111 00010111 11111111	0
11001000 00010111 00011000 00000000 through 11001000 00010111 00011000 11111111	1
11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111	2
otherwise	3

Destination-based forwarding table

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise	3

DA: 11001000 00010111 00010110 10100001

which interface?

DA: 11001000 00010111 00011000 10101010

which interface?


Longest prefix matching

- When looking for forwarding table entry for a given destination address, use the longest address prefix that matches destination address.

Question!



How many prefixes are there?



How large is the forwarding table?

```
rviews@route-server.ip.att.net> show route summary
Autonomous system number: 65000
Router ID: 12.0.1.28

inet.0: 763336 destinations, 12211712 routes (763336 active, 0 holddown, 0 hidden)
  Direct:      1 routes,      1 active
  Local:       1 routes,      1 active
  BGP: 12211603 routes, 763227 active
  Static:     107 routes,    107 active

inet6.0: 73069 destinations, 1168998 routes (73069 active, 0 holddown, 0 hidden)
  Direct:      1 routes,      1 active
  Local:       2 routes,      2 active
  BGP: 1168992 routes, 73063 active
  Static:      2 routes,      2 active
  INET6:       1 routes,      1 active
```

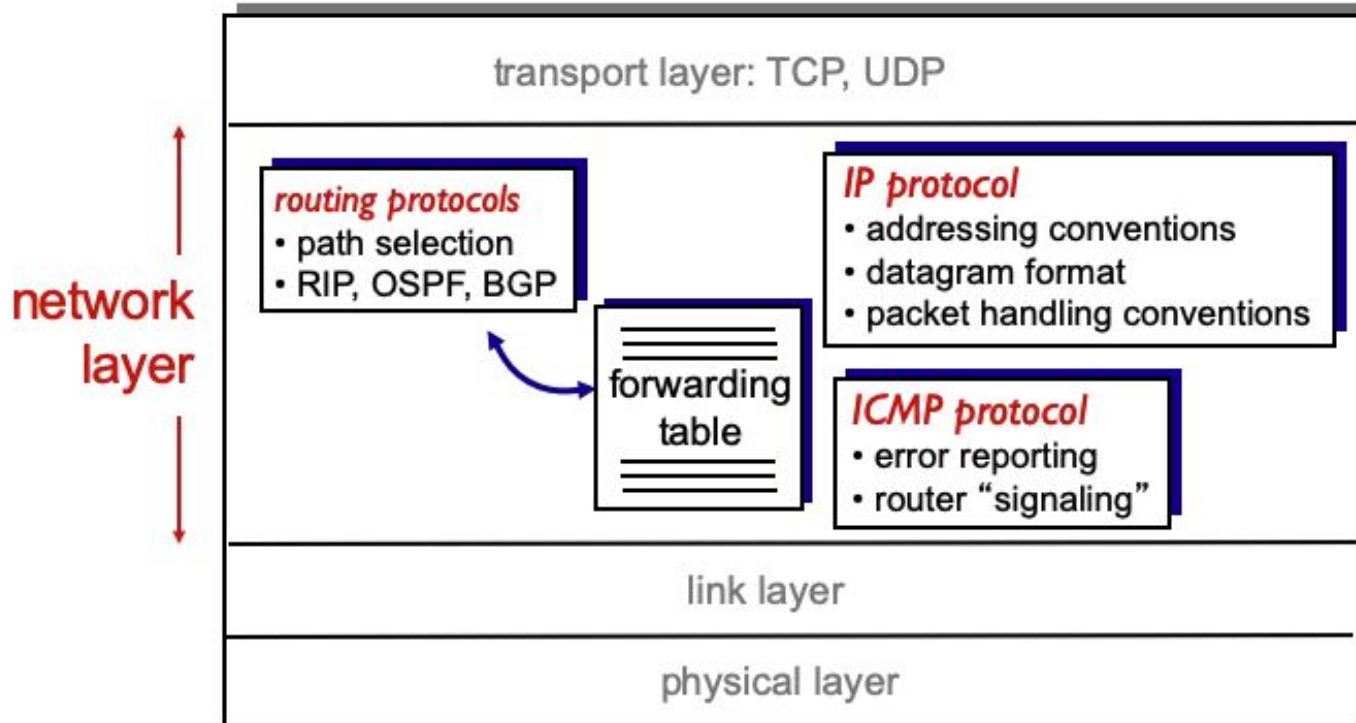
Forwarding table example

```
rviews@route-server.ip.att.net> show route forwarding-table
Routing table: default.inet
Internet:
Enabled protocols: Bridging,

```

Destination	Type	RtRef	Next hop	Type	Index	NhRef	Netif
default	user	7	0:0:5e:0:1:1	ucst	578	32	em0.0
default	perm	0		rjct	36	1	
0.0.0.0/32	perm	0		dscd	34	106	
1.93.23.118/32	user	0		dscd	34	106	
12.0.0.0/8	user	0		indr	1048574	300	
			0:0:5e:0:1:1	ucst	578	32	em0.0
12.0.0.0/9	user	0		indr	1048574	300	
			0:0:5e:0:1:1	ucst	578	32	em0.0
12.0.1.0/24	intf	0		rslv	565	1	em0.0
12.0.1.0/32	dest	0	12.0.1.0	recv	563	1	em0.0
12.0.1.1/32	dest	0	0:0:5e:0:1:1	ucst	578	32	em0.0
12.0.1.25/32	dest	1	0:6:5b:f0:d6:44	ucst	584	2	em0.0
12.0.1.28/32	intf	0	12.0.1.28	locl	564	2	
12.0.1.28/32	dest	0	12.0.1.28	locl	564	2	
12.0.1.62/32	dest	0	0:13:80:d1:64:40	ucst	593	1	em0.0
12.0.1.71/32	dest	1	0:21:5e:c8:a4:78	ucst	580	2	em0.0
12.0.1.139/32	dest	1	0:5:85:ce:1d:f4	ucst	579	2	em0.0
12.0.1.148/32	dest	1	52:54:0:42:2c:ed	ucst	585	2	em0.0
12.0.1.160/32	dest	1	52:54:0:1:8e:4	ucst	586	2	em0.0

Network layer

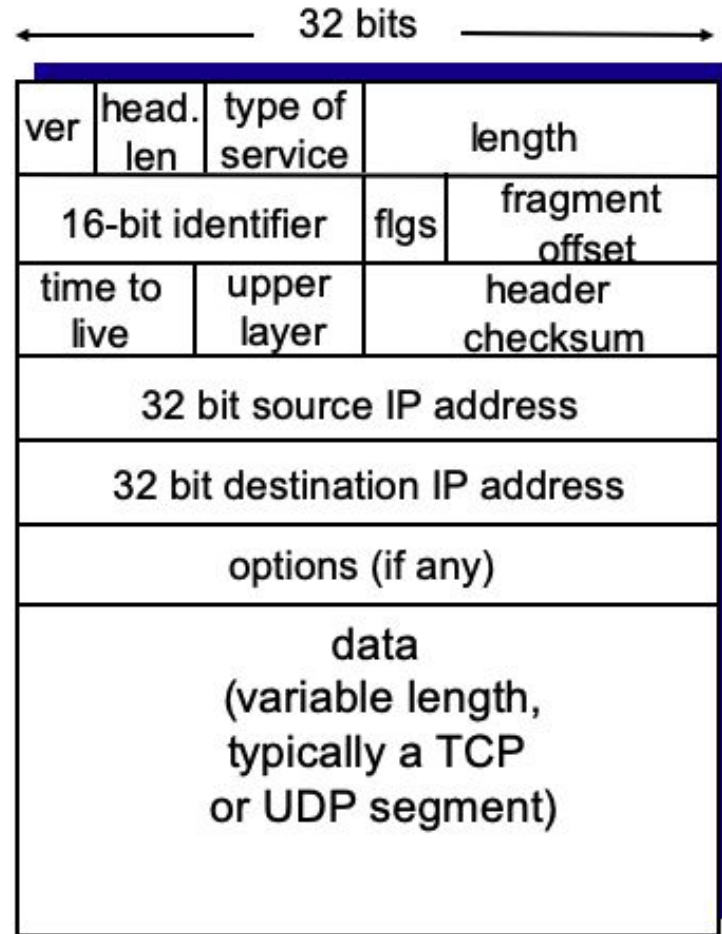


IP (v4 and v6)

IPv4 datagram format


how much overhead?

- ❖ 20 bytes of TCP
- ❖ 20 bytes of IP
- ❖ = 40 bytes + app layer overhead

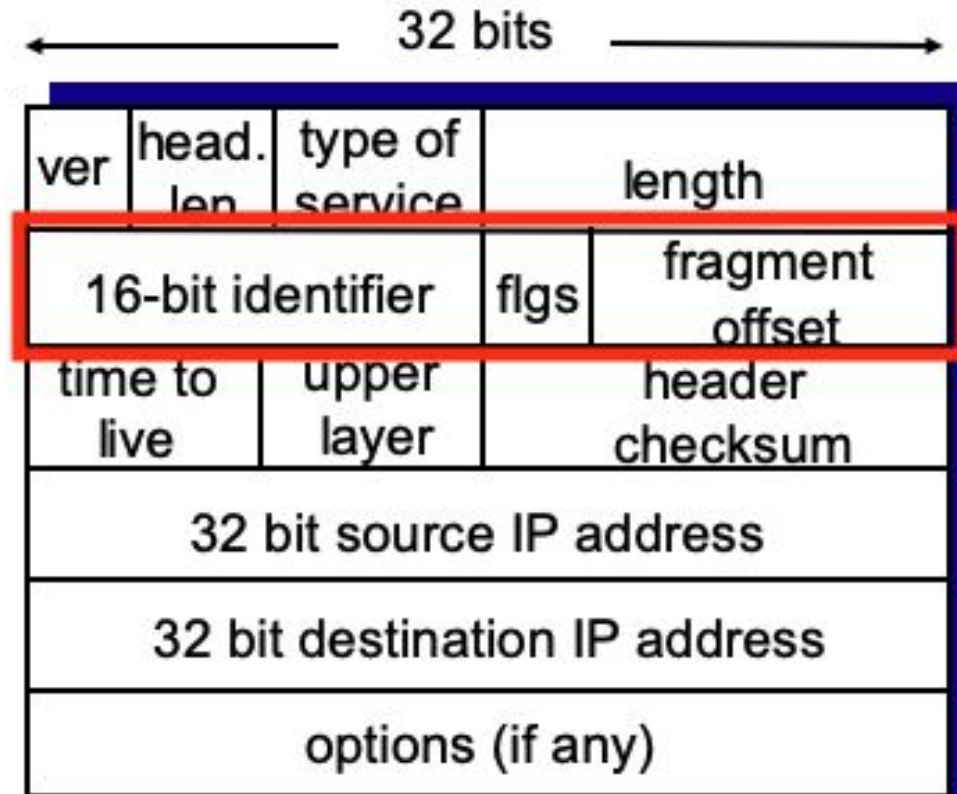


IP fragmentation, reassembly

- Network links have MTU (max transfer size), aka largest possible link level frame
- Large IP datagram will be “fragmented”
- IP header bits used to identify, order related fragments



IP fragmentation, assembly



IP fragmentation, assembly

example:

- ❖ 4000 byte datagram
- ❖ MTU = 1500 bytes

	length	ID	fragflag	offset
	=4000	=x	=0	=0

*one large datagram becomes
several smaller datagrams*

1480 bytes in
data field

offset =
 $1480/8$


	length	ID	fragflag	offset
	=1500	=x	=1	=0

	length	ID	fragflag	offset
	=1500	=x	=1	=185


	length	ID	fragflag	offset
	=1040	=x	=0	=370

Why?

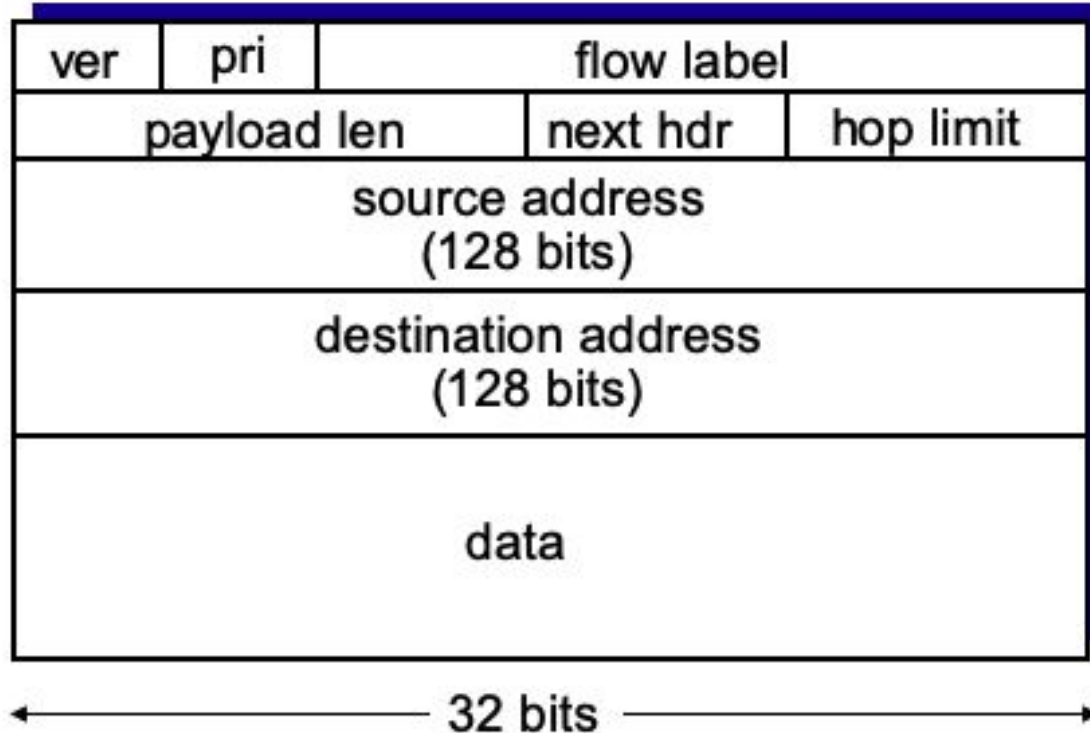
- Send packet from A to B
- Probability of a success delivery is p
- What happens in case of fragmentation?



IPv6 motivation



IPv6 datagram format



IPv6 header

- Fixed length 40 bytes
 - 32 bytes are used for source and destination IP addresses
- No checksum
 - Lower layer protocols have CRC to detect errors
- Hop limit is the same as TTL in v4

How many addresses?

- Land 148,940,000 km²
- Water 361,132,000 km²
- Total 510,072,000 km²
- Number of IPv6 addresses 2^{128}
- 66,712,614,478,140,039,732,307 IP addresses per ft²

IPv6 address notation

- 128 bits noted as eight 16-bit fields
- Separated by colons, not dots
- Each integer is represented by 4 hexadecimal digits
- E.g. 2001:0DB8:0000:0000:0000:0000:3257:9652

IPv6 shorthand


- It is likely that at first there may be many many zero's in the address.
- FF01:0000:0000:0000:0000:0000:0000:0001
- FF01:0:0:0:0:0:0:1
- FF01:0:0::1
- FF01::1
- 0:0:0:0:0:0:0:1 = ::1
- 0:0:0:0:0:0:0:0 = ::

IPv6 network notation

- IPv6 network address are denoted by CIDR notation
- The initial bits of IPv6 address form the network prefix
- E.g. 2001:CDBA:9ABC:5678::/64
 - 2001:CDBA:9ABC:5678:: to
2001:CDBA:9ABC:5678::FFFF:FFFF:FFFF:FFFF

IPv6 network notation

- Network prefix 64 bits, host identifier 64 bits
- ISP allocates you /48
- 2^{16} gives 65536 subnets
- You allocate a /64 to each interface




Routing protocols

Routing

- Determine route taken by packets from source to destination
- Two approaches
 - per-router control (traditional)
 - Logically centralized control (software defined networking, aka SDN)


Per-router control

Individual routing algorithm components in each and every router, interact with each other to compute forwarding tables.



Logically centralized control

A distinct remote controller interacts with local control agents (CAs) in routers to compute forwarding tables




Goal of routing protocols

- Determine “good” paths from source to destination
- Path is a sequence of routers packets will traverse



Graph abstraction of the network




graph: $G = (N, E)$

N = set of routers = $\{ u, v, w, x, y, z \}$

E = set of links = $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

Graph abstraction of the network



$c(x, x') = \text{cost of link } (x, x')$
e.g., $c(w, z) = 5$

cost could always be 1, or
inversely related to bandwidth,
or inversely related to
congestion

cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

key question: what is the least-cost path between u and z ?

routing algorithm: algorithm that finds that least cost path

Routing algorithm classification

- Global
 - All routers have complete topology, link cost info
 - **“Link state” algorithm**
- Decentralized
 - Router only knows physically connected neighbors, link costs to neighbors
 - Iterative process of computation and exchange info
 - **“Distance vector” algorithm**

Link-state routing algorithm

- Use Dijkstra's algorithm
- Network topology, link cost known to all nodes
 - via link-state broadcast, all nodes have same info
- Compute least cost paths from one to all other nodes
 - produce forwarding table for that node

Distance vector algorithm

- Distributed version of Bellman-Ford

let

$d_x(y) :=$ cost of least-cost path from x to y

then

$$d_x(y) = \min_v \{ c(x,v) + d_v(y) \}$$


cost from neighbor v to destination y

cost to neighbor v

\min taken over all neighbors v of x

Distance vector algorithm

clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$



$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

Distance vector algorithm

- $D_x(y)$ = estimate of least cost from x to y
- x maintains distance vector $D_x = [D_x(y) \text{ for } y \text{ in } N]$
- Node x :
 - knows cost to each neighbor v which is $c(x,v)$
 - maintains its neighbors distance vectors. For each neighbor v , x maintains $D_v = [D_v(y) \text{ for all } y \text{ in } N]$

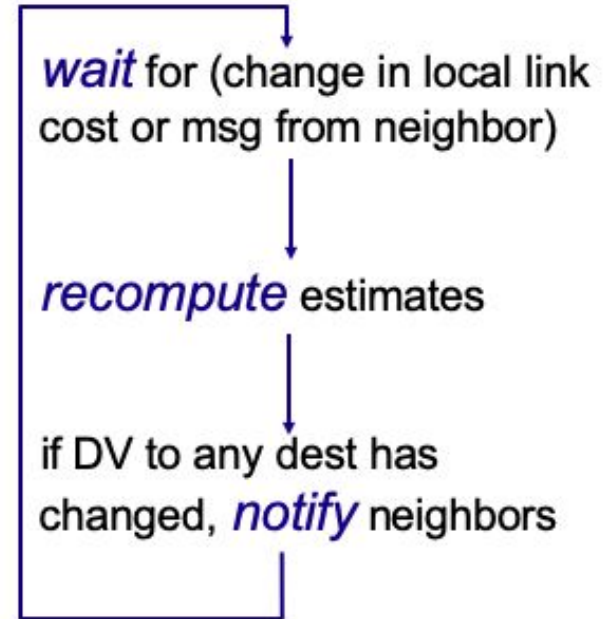
Distance vector algorithm

- Key idea
 - from time to time, each node sends its own distance vector estimate to neighbors
 - when x receives new distance vector from neighbor, it updates its own distance vector

$$D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\} \text{ for each node } y \in N$$

Distance vector algorithm

- Iterative and asynchronous. Each iteration is caused by
 - local link cost change
 - distance vector update from neighbors
- Distributed
 - only notify its own neighbors



**node x
table**

	cost to		
	x	y	z
from x	0	2	7
from y	∞	∞	∞
from z	∞	∞	∞

cost to


	x	y	z
from x	0		
from y	2	0	1
from z	7	1	0

**node y
table**

	cost to		
	x	y	z
from x	∞	∞	∞
from y	2	0	1
from z	∞	∞	∞

**node z
table**

	cost to		
	x	y	z
from x	∞	∞	∞
from y	∞	∞	∞
from z	7	1	0



time

**node x
table**

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞

cost to


		x	y	z
from	x	0	2	
	y	2	0	1
	z	7	1	0

**node y
table**

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞

**node z
table**

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0



time

**node x
table**

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞


		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0

**node y
table**

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞

**node z
table**

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0



time

**node x
table**

from \ cost to	x	y	z
	x	y	z
x	0	2	7
y	∞	∞	∞
z	∞	∞	∞

**node y
table**

from \ cost to	x	y	z
	x	y	z
x	∞	∞	∞
y	2	0	1
z	∞	∞	∞

**node z
table**

from \ cost to	x	y	z
	x	y	z
x	∞	∞	∞
y	∞	∞	∞
z	7	1	0

from \ cost to	x	y	z
	x	y	z
x	0	2	3
y	2	0	1
z	7	1	0


from \ cost to	x	y	z
	x	y	z
x	0	2	7
y	2	0	1
z	7	1	0

from \ cost to	x	y	z
	x	y	z
x	0	2	7
y	2	0	1
z	3	1	0

from \ cost to	x	y	z
	x	y	z
x	0	2	3
y	2	0	1
z	3	1	0

from \ cost to	x	y	z
	x	y	z
x	0	2	3
y	2	0	1
z	3	1	0

from \ cost to	x	y	z
	x	y	z
x	0	2	3
y	2	0	1
z	3	1	0





time →

Link cost change

- Node detects local link cost change
- Update routing info, recalculates distance vector
- Notify neighbors if distance vector changes



What happens when cost reduces?



Distance vector algorithm



What happens when cost increases



Distance vector



BAD NEWS TRAVELS SLOW

Poisoned reverse

- If z routes through y to get to x
 - z tells y its distance to x is infinite
 - so y won't route to x via z



Comparison between LS and DV

- Message complexity
 - LS: with n nodes and E links, $O(nE)$ msg sent
 - DV: it varies
- Speed of convergence
 - LS: with n nodes and E links, $O(n^2)$
 - DV: it varies

PICK ONE ALGORITHM

**APPLY TO THE
WHOLE INTERNET**


No!!!!!!

- Scalability
 - neither can scale to handle the entire Internet
- Administrative autonomy
 - Internet = network of networks
 - each network admin may want to control routing in its own network

The Internet approach

- Aggregate routers into regions known as “autonomous systems” (aka AS)
- intra-AS routing
 - routing among hosts/routers in the same AS
 - routers in same AS run the same protocol
 - routers in different AS can run different protocols
- inter-AS routing
 - routing among AS'es

The Internet approach



Intra-AS routing

- Also known as Interior Gateway Protocols (IGP)
- Most common intra-AS routing protocols
 - RIP: Routing Information Protocol
 - based on distance vector
 - OSPF: Open Shortest Path First
 - based on link state
 - EIGRP: Enhanced Interior Gateway Routing Protocol
 - Proprietary protocol by Cisco

Inter-AS routing


- BGP: Border Gateway Protocol
 - the de facto inter-AS routing protocol
- No other inter-AS routing protocols!



BGP

- BGP provides each AS a means to
 - eBGP: obtain subnet reachability information from neighboring AS'es
 - iBGP: propagate reachability information to all AS-internal routers
 - Determine “good” routes to other networks based on reachability information and policy

BGP



BGP basics

- Designed to scale huge inter network like the Internet
- Send updates to manually defined neighbors
- Application layer protocol using TCP (port 179)
- AS-path information to prevent loop
- Path selection is complicated

BGP neighbors

- BGP neighbors are routers forming TCP connections to exchange BGP updates
 - manually configured
 - two types of neighbor relationship
 - iBGP (routers in the same AS)
 - eBGP (routers in different AS)

BGP neighbors

```
route-server> show ip bgp neighbors
BGP neighbor is 216.218.252.130, remote AS 6939, local AS 6939, internal link
  BGP version 4, remote router ID 216.218.252.130
  BGP state = Established, up for 00:41:47
  Last read 00:00:00, hold time is 180, keepalive interval is 60 seconds
  Neighbor capabilities:
    4 Byte AS: advertised and received
    Route refresh: advertised and received(old & new)
    Address family IPv4 Unicast: advertised and received
    Graceful Restart Capabilty: advertised
  Message statistics:
```

BGP neighbors

```
route-server> show ip bgp summary
BGP router identifier 64.62.142.154, local AS number 6939
RIB entries 1449658, using 155 MiB of memory
Peers 46, using 408 KiB of memory
```


Neighbor	V	AS	MsgRcvd	MsgSent	TblVer	InQ	OutQ	Up/Down	State/PfxRcd
216.218.252.130	4	6939	3363817	6381	0	0	0	01:37:41	795891
216.218.252.147	4	6939	3645934	6379	0	0	0	01:37:39	796243
216.218.252.151	4	6939	3632194	6379	0	0	0	01:37:45	796245
216.218.252.154	4	6939	3260700	6382	0	0	0	01:37:25	795890
216.218.252.157	4	6939	3649388	6381	0	0	0	01:37:18	796245
216.218.252.164	4	6939	3505378	6379	0	0	0	01:38:15	796247
216.218.252.165	4	6939	3680065	6378	0	0	0	01:37:59	796247
216.218.252.167	4	6939	3453321	6382	0	0	0	01:37:57	795869
216.218.252.168	4	6939	3617330	6383	0	0	0	01:37:23	796244
216.218.252.169	4	6939	3123208	6378	0	0	0	01:37:32	796245
216.218.252.171	4	6939	3606008	6385	0	0	0	01:37:34	796243
216.218.252.173	4	6939	3243682	6382	0	0	0	01:37:21	795892
216.218.252.174	4	6939	3107783	6386	0	0	0	01:37:54	795888
216.218.252.176	4	6939	3363895	6382	0	0	0	01:37:16	796247
216.218.252.177	4	6939	3437999	6380	0	0	0	01:37:48	795869
216.218.252.178	4	6939	3526745	6382	0	0	0	01:37:30	796180

BGP route advertisement

- BGP routers send updates to neighbors
 - list of network prefixes is not enough
 - paths to different destination network prefixes
 - AS-PATH attribute
 - BGP is a “path vector” protocol

BGP route advertisement

- When AS3 gateway router 3a advertises path “AS3,X” to AS2 gateway router 2c, AS3 promises to AS2 it will forward datagrams towards network X



BGP attributes

- Advertised prefix includes BGP attributes
 - prefix + attributes = “route”
- Two important attributes
 - AS-PATH: list of AS'es through which prefix advertisement has passed
 - NEXT-HOP: indicates specific internal-AS router to next hop AS

BGP attributes

```
route-server> show ip bgp
```

```
BGP table version is 0, local router ID is 64.62.142.154
```

```
Status codes: s suppressed, d damped, h history, * valid, > best, = multipath,  
               i internal, r RIB-failure, S Stale, R Removed
```

```
Origin codes: i - IGP, e - EGP, ? - incomplete
```


Network	Next Hop	Metric	LocPrf	Weight	Path
* i1.0.0.0/24	216.218.252.168	1	100	0	13335 i
* i	216.218.252.173	1	100	0	13335 i
* i	198.32.146.195	1	100	0	13335 i
* i	216.218.252.179	1	100	0	13335 i
* i	216.218.252.169	1	100	0	13335 i
* i	216.218.252.184	1	100	0	13335 i
*>i	216.218.252.130	1	100	0	13335 i
* i1.0.5.0/24	216.218.252.180	1	140	0	4826 38803 56203 i
* i	64.71.184.46	100	140	0	4826 38803 56203 i

Policy-based routing

- Gateway receiving route advertisement uses import policy to accept/decline path
 - e.g. never route through AS Y
- AS policy also determines whether to advertise path to other neighboring AS'es


BGP advertisement

- AS2 router 2c receives path advertisement "AS3,X" (via eBGP) from AS3 router 3a




BGP advertisement

- Based on AS2 policy, AS2 router 2c accepts path “AS3,X” and propagates (via iBGP) to AS2 routers




BGP advertisement

- Based on AS2 policy, AS2 router 2a advertises (via eBGP) path "AS2,AS3,X" to AS1 router 1c



BGP advertisement

- Gateway router may learn about multiple paths to destination




BGP route selection

- Router may learn about more than one route to destination prefix, select route based on
 - local preference value attribute
 - shortest AS-PATH
 - closest NEXT-HOP router
 - additional criteria


Hot potato routing

- 2d learns (via iBGP) it can route to X via 2a or 2c
- Which one should 2d choose?




Hot potato routing

- Hot potato routing: choose local gateway that has the least intra domain cost



BGP prefix advertisement

- A, B, C are provider networks
- X, Y, Z are customers
 - X is dual-homed




legend:

 provider network

 customer network:

BGP prefix advertisement

- X needs to be careful when designing policy
- X does not want to route from B to C via X
 - so X will not advertise to B a route to C



So far...


- Main functionality in network layer
 - data plane: forwarding
 - control plane: routing
- Routing algorithms
 - intra-domain
 - link-state (e.g. OSPF), distance vector (e.g. RIP)
 - inter-domain
 - BGP

SDN - software defined networking

- Internet network layer
 - Historically has been implemented via distributed, per-router approach
 - Monolithic router contains switching hardware, runs proprietary implementation of standard protocols (e.g. OSPF, BGP) in proprietary router OS

SDN

- Individual routing algorithm components in each and every router interact with each other in control plane




That implies

- Limitation in scalability
- DevOps headache
- Harder to make changes
- Incompatibility between different vendors
- etc. etc.

SDN



- Logically centralized control plane




SDN - goals

- Easier network management
- Programmable
 - centralized “programming” is easier
 - compute forwarding table centrally and distribute
 - distributed “programming” is much harder
- Open (non-proprietary)

Analogy

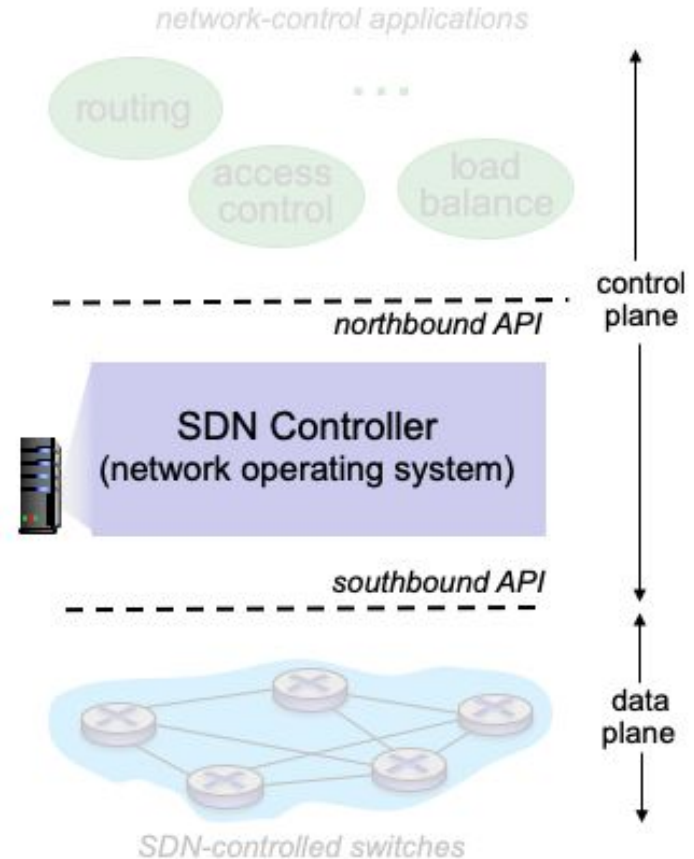


Analogy



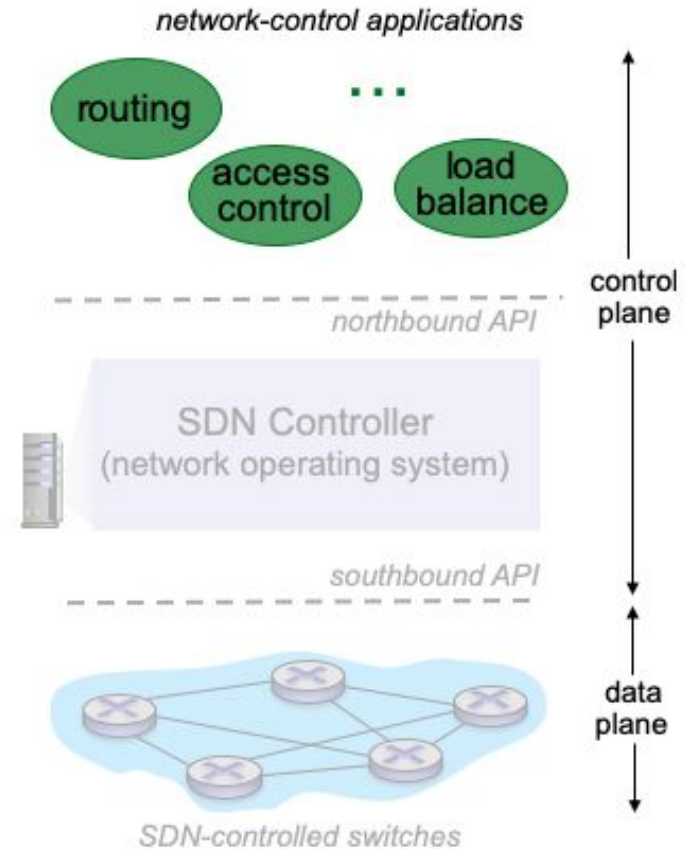
SDN - controller

- Maintain network state information
- Interact with network devices “below” via southbound API
- Interact with application “above” via northbound API
- Implemented as distributed system for scalability, availability



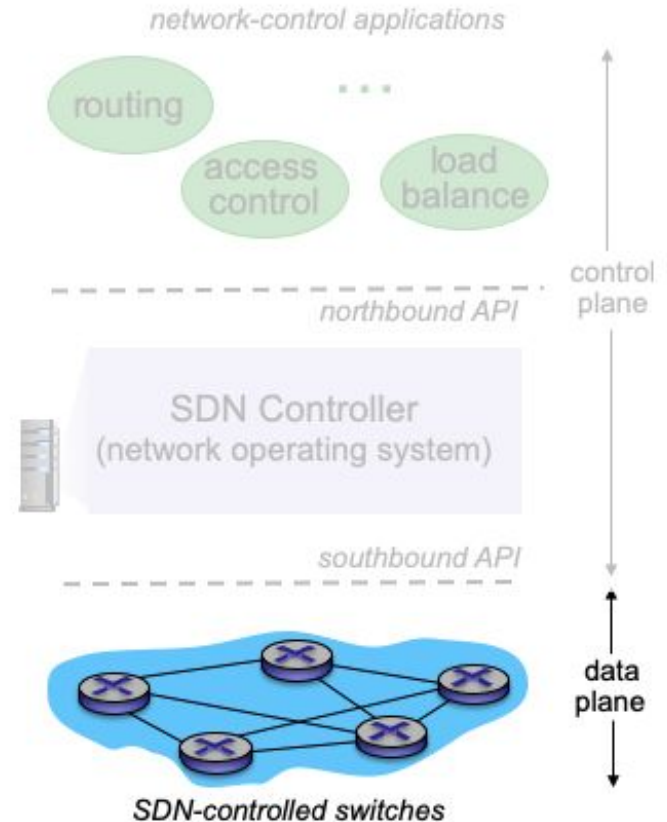
SDN - application

- Implement control functions using lower-level services provided by SDN controller



SDN - data plane

- Fast, simple, commodity layer 2 & 3 devices implementing generalized data-plane forwarding
- Table computed and installed by controller
- Protocol for communicating with controller (e.g. OpenFlow)



ICMP

ICMP

- Internet Control Message Protocol
- Used by hosts & routers to communicate network level information
 - e.g. echo request/reply (used by ping)
- Network layer “above” IP
 - carried in IP datagram


ICMP

- ICMP message
 - type/code
 - checksum
 - other header fields
 - depends on (type/code)
 - partial original IP datagram

<u>Type</u>	<u>Code</u>	<u>description</u>
0	0	echo reply (ping)
3	0	dest. network unreachable
3	1	dest host unreachable
3	2	dest protocol unreachable
3	3	dest port unreachable
3	6	dest network unknown
3	7	dest host unknown
4	0	source quench (congestion control - not used)
8	0	echo request (ping)
9	0	route advertisement
10	0	router discovery
11	0	TTL expired
12	0	bad IP header

Traceroute

- Source sends series of UDP segment to destination
 - Set TTL=1, TTL=2, etc.
- When datagram with TTL=n arrives to nth router
 - router discards datagram and sends source ICMP message (type 11, code 0)
 - ICMP message include name of router & IP address



Summary

- Forwarding & routing
- IP v4 and v6
- Routing protocols
 - Intra-domain & inter-domain
- ICMP

Questions?

